

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

df = pd.read_csv("Ecommerce Customers.txt")
df.head()
```

	Email \
0	mstephenson@fernandez.com
1	hduke@hotmail.com
2	pallen@yahoo.com
3	riverarebecca@gmail.com
4	mstephens@davidson-herman.com

	Address	Avatar
0	835 Frank Tunnel\nWrightmouth, MI 82180-9605	Violet
1	4547 Archer Common\nDiazchester, CA 06566-8576	DarkGreen
2	24645 Valerie Unions Suite 582\nCobbborough, D...	Bisque
3	1414 David Throughway\nPort Jason, OH 22070-1220	SaddleBrown
4	14023 Rodriguez Passage\nPort Jacobville, PR 3...	MediumAquaMarine

	Avg. Session Length	Time on App	Time on Website	Length of Membership \
0	34.497268	12.655651	39.577668	4.082621
1	31.926272	11.109461	37.268959	2.664034
2	33.000915	11.330278	37.110597	4.104543
3	34.305557	13.717514	36.721283	3.120179
4	33.330673	12.795189	37.536653	4.446308

	Yearly Amount Spent
0	587.951054
1	392.204933
2	487.547505
3	581.852344
4	599.406092

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 500 entries, 0 to 499
Data columns (total 8 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Email                                500 non-null    object
1   Address                             500 non-null    object
2   Avatar                              500 non-null    object
3   Avg. Session Length                 500 non-null    float64
4   Time on App                         500 non-null    float64
5   Time on Website                     500 non-null    float64
6   Length of Membership                 500 non-null    float64
7   Yearly Amount Spent                 500 non-null    float64
dtypes: float64(5), object(3)
memory usage: 31.4+ KB

df.duplicated().sum()

0

df.isnull().sum()

Email                                0
Address                             0
Avatar                              0
Avg. Session Length                 0
Time on App                         0
Time on Website                     0
Length of Membership                 0
Yearly Amount Spent                 0
dtype: int64

df.shape

(500, 8)

df1 = df.drop(['Email', 'Address', 'Avatar'], axis = 'columns')
df1.head()

   Avg. Session Length  Time on App  Time on Website  Length of
Membership \
0          34.497268    12.655651         39.577668
4.082621
1          31.926272    11.109461         37.268959
2.664034
2          33.000915    11.330278         37.110597
4.104543
3          34.305557    13.717514         36.721283
3.120179
4          33.330673    12.795189         37.536653
4.446308

```

	Yearly Amount Spent
0	587.951054
1	392.204933
2	487.547505
3	581.852344
4	599.406092

```
df1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 500 entries, 0 to 499
```

```
Data columns (total 5 columns):
```

#	Column	Non-Null Count	Dtype
0	Avg. Session Length	500 non-null	float64
1	Time on App	500 non-null	float64
2	Time on Website	500 non-null	float64
3	Length of Membership	500 non-null	float64
4	Yearly Amount Spent	500 non-null	float64

```
dtypes: float64(5)
```

```
memory usage: 19.7 KB
```

```
x = df1.drop(['Yearly Amount Spent'],axis=1)
```

```
y = df1[['Yearly Amount Spent']]
```

```
from sklearn.model_selection import train_test_split
```

```
x_train,x_test,y_train,y_test =
```

```
train_test_split(x,y,test_size=0.3,random_state=101)
```

```
x_train
```

	Avg. Session Length	Time on App	Time on Website	Length of Membership
202	31.525752	11.340036	37.039514	3.811248
428	31.862741	14.039867	37.022269	3.738225
392	33.258238	11.514949	37.128039	4.662845
86	33.877779	12.517666	37.151921	2.669942
443	33.025020	12.504220	37.645839	4.051382
...
63	32.789773	11.670066	37.408748	3.414688
326	33.217188	10.999684	38.442767	4.243813
337	31.827979	12.461147	37.428997	

```

2.974737
11          33.879361    11.584783    37.087926
3.713209
351         32.189845    11.386776    38.197483
4.808320

```

```
[350 rows x 4 columns]
```

```
x_test
```

	Avg. Session Length	Time on App	Time on Website	Length of Membership
18	32.187812	14.715388	38.244115	1.516576
361	32.077590	10.347877	39.045156	3.434560
104	31.389585	10.994224	38.074452	3.428860
4	33.330673	12.795189	37.536653	4.446308
156	32.294642	12.443048	37.327848	5.084861
..
147	32.255901	10.480507	37.338670	4.514122
346	32.765665	12.506548	35.823467	3.126509
423	33.128693	10.398458	36.683393	3.859818
17	32.338899	12.013195	38.385137	2.420806
259	32.096109	10.804891	37.372762	2.699562

```
[150 rows x 4 columns]
```

```
y_train
```

	Yearly Amount Spent
202	443.965627
428	556.298141
392	549.131573
86	487.379306
443	561.516532
..	...
63	483.159721
326	505.230068
337	440.002748
11	522.337405

```
351          533.396554
```

```
[350 rows x 1 columns]
```

```
y_test
```

	Yearly Amount Spent
18	452.315675
361	401.033135
104	410.069611
4	599.406092
156	586.155870
..	...
147	479.731938
346	488.387526
423	461.112248
17	407.704548
259	375.398455

```
[150 rows x 1 columns]
```

```
from sklearn.linear_model import LinearRegression
```

```
lr = LinearRegression()
```

```
lr.fit(x_train,y_train)
```

```
LinearRegression()
```

```
lr.score(x_train,y_train)
```

```
0.9817562058732432
```

```
lr.score(x_test,y_test)
```

```
0.9890046246741234
```

```
pred = lr.predict(x_test)
```

```
plt.scatter(y_test,pred)
```

```
plt.xlabel('y_test')
```

```
plt.ylabel('pred')
```

```
plt.show()
```

